

# Evaluation

Kevin Brandstatter  
Weiwei Wu

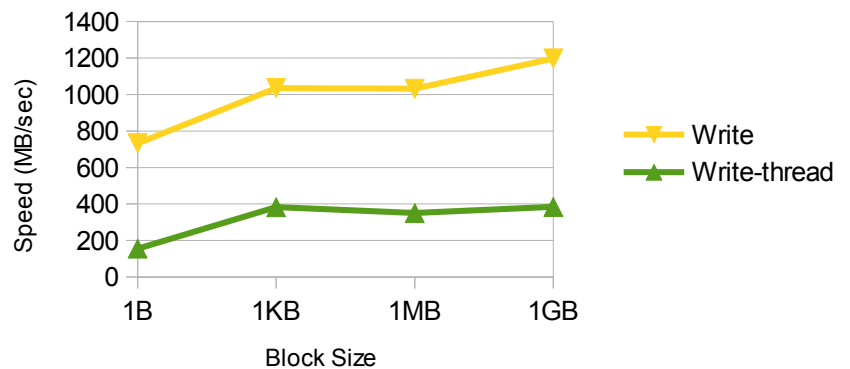
All Benchmarks done on AMD Phenom x6 @ 3.3 GHz (6 cores). Gigabit Ethernet card. Two 1 TB harddrives in Raid 1. 16 GB of RAM. System at mostly Idle for tests. Network tests done over the local loopback interface.

## Diskbench:

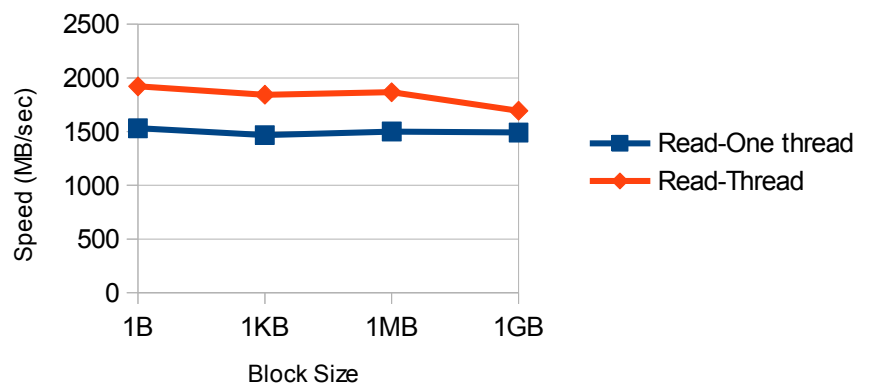
From the results of the disk benchmark its easy to conclude several things. As expected reads are much faster than writes. We can also conclude that we can get higher throughput using larger block sizes for the write operations. On the other hand, Using larger block sizes for the read operations seem to degrade performance.

On a more unexpected result, it seems that our threaded operations are much slower than their non threaded counterpart. This is likely due to the fact that the threads are competing for the read and write heads of the disk which do not parallelize very well. This causes the operations to thrash on disk thus lowering performance.

## Disk Benchmark Write




## Disk Benchmark Read



## Memory Benchmark

From our memory benchmark we can see that memcpy was the most affected by an increase in block size which indicates that the memcpy operation is very slow for large data sizes. Also, they all saw performance degrade as the number of threads increased. This is most likely due to hardware constraints of memory access.

## Memory Benchmark




| blocksize | Sequential | Sequential (threaded) | Random    | Random (threaded) | Memcpy      | Memcpy (threaded) |
|-----------|------------|-----------------------|-----------|-------------------|-------------|-------------------|
| 1B        | 105.085998 | 184.537755            | 43.241198 | 17.85525          | 16.273393   | 21.971986         |
| 1KB       | 94.731311  | 182.74619             | 45.775756 | 12.108502         | 4089.456869 | 7780.415994       |
| 1MB       | 85.508912  | 197.477533            | 41.957996 | 3.406231          | 3928.477298 | 5883.149813       |

### Cpubench:


From the results its clear to see that as we add parallelism to the process, we increase cpu performance. While this performance is not the ideal 2x it approaches 2x which is very promising given that some overhead will be used by the system for scheduling and such. Also, as expected GFLOPS are lower than GIOPS as they are more cpu intensive.

### Cpu Benchmarks




### NetBench


#### UDP Benchmarks



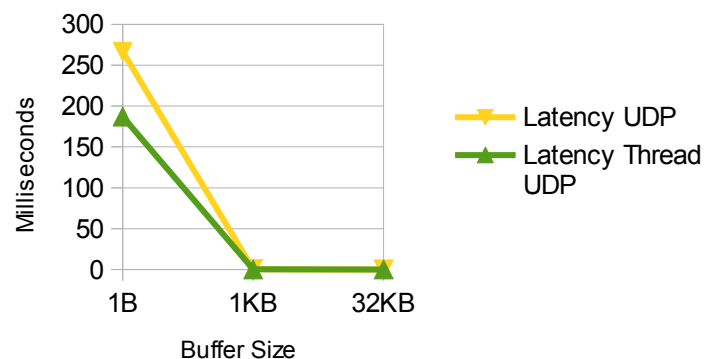
#### TCP Benchmarks



#### TCP Latency



#### Latency UDP



From the network benchmarks we can see that there is a serious benifet to threading the operations as UDP and TCP both had higher throughput and lower latency with multiple threads. Also, The

tests show that there is a benefit to using a larger buffer size. This is because with a larger buffer size, the overhead of the data transmitted is amortized into the time of sending the data itself and does not have as big of an impact as it does with smaller buffer sizes.